

Kubeadm Setup

Following are the prerequisites for Kubeadm Kubernetes cluster setup

- Minimum two **Ubuntu** nodes [One master and one worker node]. You can have more worker nodes as per your requirement.
- The master node should have a minimum of **2 vCPU and 2GB RAM**.
- For the worker nodes, a minimum of **1vCPU and 2 GB RAM** is recommended.
- 10.X.X.X network range with static IPs for master and worker nodes. We will be using the 192 series as the pod network range that will be used by the Calico network plugin. Make sure the Node IP range and pod IP range don't overlap.

Control-plane node(s)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	6443*	Kubernetes API server	All
TCP	Inbound	2379-2380	etcd server client API	kube-apiserver, etcd
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	10251	kube-scheduler	Self
TCP	Inbound	10252	kube-controller-manager	Self

Worker node(s)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	30000-32767	NodePort Services**	All

- Update the apt package index and install packages needed to use the Kubernetes apt repository:

```
$ sudo apt-get update -y
```

```
$ sudo apt-get install -y apt-transport-https ca-certificates curl
```

- Execute the following commands for IPtables to see bridged traffic.

```
ubuntu@ip-172-31-5-3:~$ cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
```

```
br_netfilter
```

```
EOF
```

```
$ cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
EOF
```

```
$ sudo systemctl --system
```

3. Disable swap on all the Nodes

For kubeadm to work properly, you need to disable swap on all the nodes using the following command.

```
$ sudo swapoff -a
```


```
$ sudo sed -i ' / swap / s/^(.*)$/#\1/g' /etc/fstab
```

4. Install Docker Container Runtime On All The Nodes

The basic requirement for a Kubernetes cluster is a container runtime. You can have any one of the following container runtimes.

 containerd

 CRI-O

 Docker

We will be using Docker for this setup.

As a first step, we need to install Docker on all the nodes. Execute the following commands on all the nodes.

Install the required packages for Docker.

```
$ sudo apt-get update -y
```

```
$ sudo apt-get install -y \
```

```
apt-transport-https \
```

```
ca-certificates \
```

```
curl \
```

```
gnupg \
```

```
lsb-release
```

5. Add the Docker GPG key and apt repository.

```
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-archive-keyring.gpg
```

```
$ echo \
```

```
"deb [arch=amd64 signed-by=/usr/share/keyrings/docker-archive-keyring.gpg]
```

```
https://download.docker.com/linux/ubuntu \
```

```
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

6. Install the Docker community edition.

```
$ sudo apt-get update -y
```

```
$ sudo apt-get install docker-ce docker-ce-cli containerd.io -y
```

7. Add the docker daemon configurations to use systemd as the cgroup driver.

```
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

Start and enable the docker service.

```
$ sudo systemctl enable docker
```

```
$ sudo systemctl daemon-reload
```

```
$ sudo systemctl restart docker
```

8. Install Kubeadm & Kubelet & Kubectl on all Nodes

Install the required dependencies.

```
$ sudo apt-get update
$ sudo curl -fsSL /usr/share/keyrings/kubernetes-archive-keyring.gpg
https://packages.cloud.google.com/apt/doc/apt-key.gpg
```

9. Add the GPG key and apt repository.

```
$ echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

10. Update apt and install kubelet, kubeadm, and kubectl.

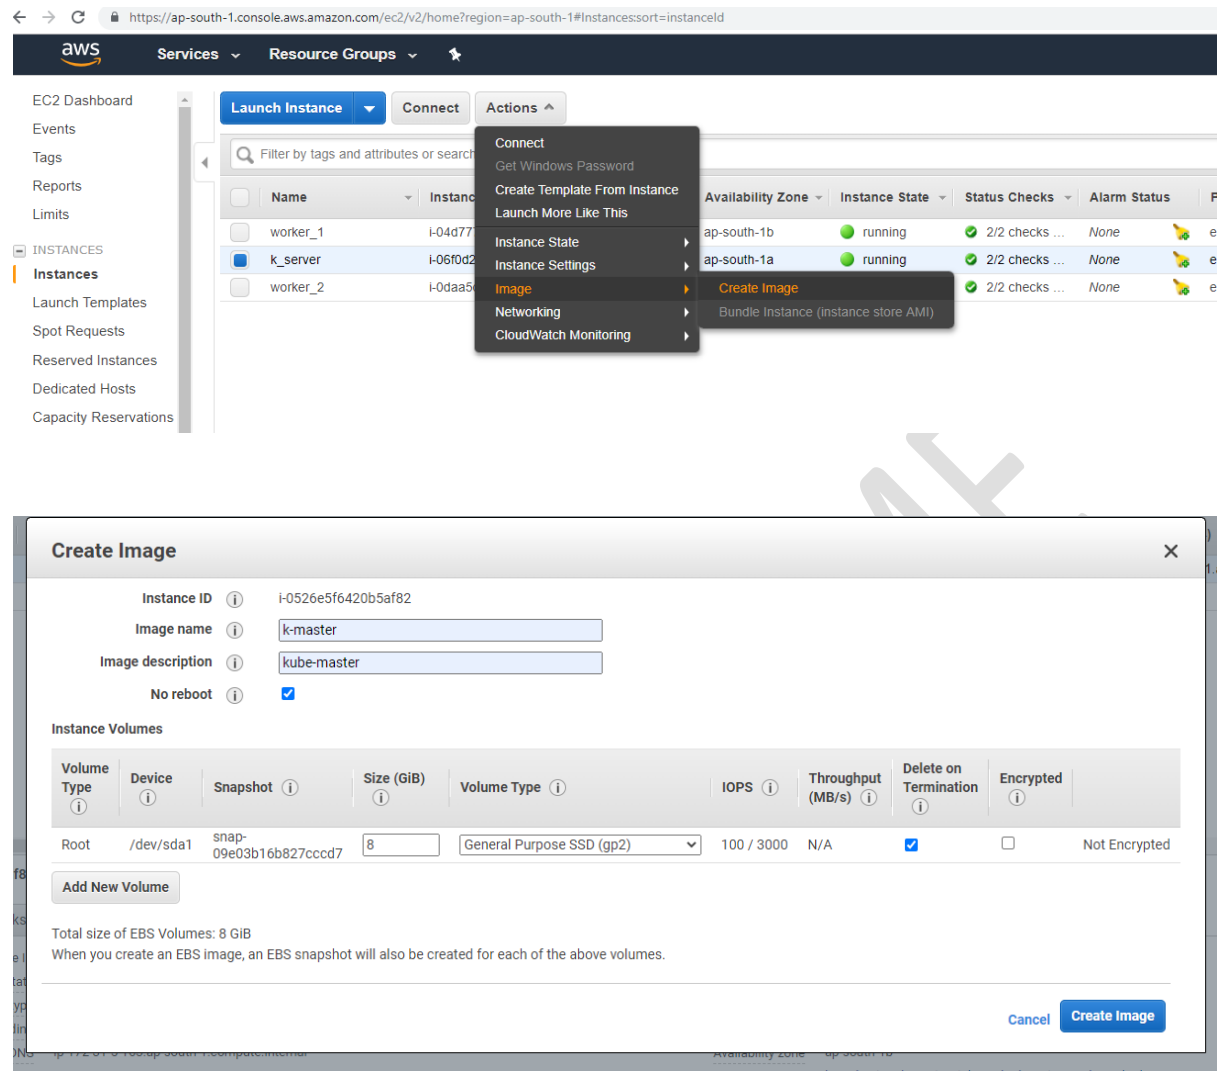
```
$ sudo apt-get update -y
$ sudo apt-get install -y kubelet kubeadm kubectl
```

```
$ sudo rm /etc/containerd/config.toml
$ sudo systemctl restart containerd
```

Now we have all the required utilities and tools for configuring Kubernetes components using kubeadm.

11. create a image After successfully install we need make it as a AMI image.

Goto aws console select which server you install all the above steps goto ACTION -> IMAGE -> CREATE IMAGE



12. Add hold to the packages to prevent upgrades.
\$ sudo apt-mark hold kubelet kubeadm kubectl

13. Initialize Kubeadm On Master Node To Setup Control Plane
Execute the commands in this section only on the master node.

First, set two environment variables. Replace 172.31.8.227 with the IP of your master node.

```
$ export IPADDR="172.31.12.191"
$ export NODENAME=$(hostname -s)
```

14. Configuring a cgroup driver
<https://kubernetes.io/docs/tasks/administer-cluster/kubeadm/configure-cgroup-driver/>

```
$ vi config.yaml
apiVersion: kubelet.config.k8s.io/v1beta1
```

```
kind: KubeletConfiguration
```

```
cgroupDriver: containerd
```

```
---
```

```
apiVersion: kubeadm.k8s.io/v1beta3
```

```
kind: ClusterConfiguration
```

```
networking:
```

```
  podSubnet: "10.244.0.0/16"
```

15. Now, initialize the master node control plane configurations using the following kubeadm command.

```
$ sudo kubeadm init --config config.yaml
```

If you get below error

```
ubuntu@ip-172-31-12-191:~$ sudo kubeadm init --config config.yaml
W0612 17:35:49.254652 5783 common.go:83] your configuration file uses a deprecated API spec: "kubeadm.k8s.io/v1beta2". Please use 'kubeadm config migrate --old-config old.yaml --new-config new.yaml', which will write the new, similar spec using a newer API version.
[init] Using Kubernetes version: v1.24.1
[preflight] Running pre-flight checks
[W0612 17:35:49.254652 5783 common.go:83] [WARNING SystemVerification]: missing optional cgroups: blkio
error execution phase preflight: [preflight] Some fatal errors occurred:
[ERROR CRI]: container runtime is not running: output: E0612 17:35:49.898269 5791 remote_runtime.go:925] "Status from runtime service failed" err="rpc error: code = Unimplemented desc = unknown service runtime.v1alpha2.RuntimeService"
time="2022-06-12T17:35:49Z" level=fatal msg="getting status of runtime: rpc error: code = Unimplemented desc = unknown service runtime.v1alpha2.RuntimeService"
, error: exit status 1
[preflight] If you know what you are doing, you can make a check non-fatal with `--ignore-preflight-errors=...`
To see the stack trace of this error execute with --v=5 or higher
```

Execute below commands ,

```
$ sudo rm /etc/containerd/config.toml
```

```
$ sudo systemctl restart containerd
```

And then initialize the kubeadm

```
$ sudo kubeadm init --config config.yaml
```

You will get below output

```
$ sudo kubeadm init --config config.yaml
```

```
W0612 17:40:08.635822 6116 common.go:83] your configuration file uses a deprecated API spec: "kubeadm.k8s.io/v1beta2". Please use 'kubeadm config migrate --old-config old.yaml --new-config new.yaml', which will write the new, similar spec using a newer API version.
```

```
[init] Using Kubernetes version: v1.24.1
```

```
[
```

```
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
```

[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key

[addons] Applied essential addon: CoreDNS

[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.12.191:6443 --token lpm1gt.dny774qp3f49hffr \  
--discovery-token-ca-cert-hash  
sha256:aa95ab7da30d88bd611c1d4de2c9e14a8edbac26f846d0ec801e64d61fbda25f
```

Use the following commands from the output to create the **kubeconfig** in master so that you can use **kubectl** to interact with cluster API.

```
$ mkdir -p $HOME/.kube  
$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Now, verify the **kubeconfig** by executing the following **kubectl** command to list all the pods in the **kube-system** namespace.

\$ kubectl get po -n kube-system

```
ubuntu@ip-172-31-12-191:~$ kubectl get po -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
coredns-6d4b75cb6d-5pbd1           0/1     Pending   0           3m17s
coredns-6d4b75cb6d-kftg2           0/1     Pending   0           3m17s
etcd-ip-172-31-12-191               1/1     Running   0           3m32s
kube-apiserver-ip-172-31-12-191     1/1     Running   0           3m32s
kube-controller-manager-ip-172-31-12-191 1/1     Running   0           3m32s
kube-proxy-xz1bx                    1/1     Running   0           3m17s
kube-scheduler-ip-172-31-12-191     1/1     Running   0           3m31s
ubuntu@ip-172-31-12-191:~$
```

```
ubuntu@ip-172-31-12-191:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-12-191    NotReady control-plane 5m40s  v1.24.1
ubuntu@ip-172-31-12-191:~$
```

You should see the following output. You will see the two Coredns pods in a pending state. It is the expected behavior. Once we install the network plugin, it will be in a running.

13. Installing a CNI Network.

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

A network is needed to enable the pods to communicate with each other. WEAVE CNI plugin is the network plugin used here.

switch the root user and run. (sudo su -)

```
ubuntu@ip-172-31-39-7:~$ sudo su -
root@ip-172-31-39-7:~# sysctl net.bridge.bridge-nf-call-iptables=1
switch user to Ubuntu
root@ip-172-31-39-7:~# su ubuntu
ubuntu@ip-172-31-39-7:~$ export kubever=$(kubectl version | base64 | tr -d '\n')
ubuntu@ip-172-31-39-7:~$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$kubever" (not working)
or
kubectl apply -f https://github.com/weaveworks/weave/releases/download/v2.8.1/weave-daemonset-k8s.yaml
```

Now that the CNI network has been created, give it a minute or 2 and test again. The result is as shown.

```
ubuntu@ip-172-31-39-7:~$ kubectl get nodes
```



```
ubuntu@ip-172-31-12-191:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
ip-172-31-12-191    Ready    control-plane  8m30s    v1.24.1
ubuntu@ip-172-31-12-191:~$
```

Now master node is ready and get the all pods and see the coredns pods are running or not

```
$ kubectl get po -n kube-system
```

```
$ kubectl get po -A
```

```
ubuntu@ip-172-31-12-191:~$ kubectl get po -A
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  coredns-6d4b75cb6d-5pbd1              1/1     Running   0           9m30s
kube-system  coredns-6d4b75cb6d-kftg2              1/1     Running   0           9m30s
kube-system  etcd-ip-172-31-12-191                 1/1     Running   0           9m45s
kube-system  kube-apiserver-ip-172-31-12-191        1/1     Running   0           9m45s
kube-system  kube-controller-manager-ip-172-31-12-191 1/1     Running   0           9m45s
kube-system  kube-proxy-xz1bx                      1/1     Running   0           9m30s
kube-system  kube-scheduler-ip-172-31-12-191        1/1     Running   0           9m44s
kube-system  weave-net-6dn8f                       2/2     Running   1 (85s ago)  91s
ubuntu@ip-172-31-12-191:~$
```

Here the CNI weave network has been installed and as a result the master node is now showing 'READY' and kube-dns is now showing 'Running' instead of pending. You will also notice the creation of a weave container weave-net-6dn8f listed above.


Now lets Creating the Kubernetes Slave Nodes

Create an instance using the AMI that was created above. For test purposes, its OK to choose t2.micro image.

Go to aws console launch instance and click left pane MY AMIs and select.

Step 1: Choose an Amazon Machine Image (AMI)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. You can select an AMI provided by AWS, our user community, or the AWS Marketplace; or you can select one of your own AMIs.



Got further and launch a instance.

Join the first Node to the Cluster. ssh into the slave node

run the join command from kubeadm init screen output above.

Don't forgot add the security group mention above first line document.

Edit instance security group in master to allow TCP 6783 and UDP 6783/6784 ports

If you forgot copy join command while you kubeadm init time use below command to get kubeadm join command.


```
$ kubeadm token create --print-join-command
```

```
$ sudo kubeadm join 172.31.12.191:6443 --token lpm1gt.dny774qp3f49hffr \  
--discovery-token-ca-cert-hash sha256:aa95ab7da30d88bd611c1d4de2c9e14a8edbac26f846d0ec801e64d61fbda25f
```

output look like below.

```
ubuntu@ip-172-31-14-178:~$ sudo kubeadm join 172.31.12.191:6443 --token lpm1gt.dny774qp3f49hffr --discovery-token-ca-cert-hash sha256:aa95ab7da30d88bd611c1d4de2c9e14a8edbac26f846d0ec801e64d61fbda25f  
[preflight] Running pre-flight checks  
[WARNING SystemVerification]: missing optional cgroups: blkio  
[preflight] Reading configuration from the cluster...  
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
```

```
ubuntu@ip-172-31-12-191:~$ kubectl get node  
NAME                STATUS    ROLES    AGE   VERSION  
ip-172-31-12-191    Ready     control-plane   24m   v1.24.1  
ip-172-31-14-178    Ready     <none>         42s   v1.24.1  
ubuntu@ip-172-31-12-191:~$
```

Kuberentes dash board setup

```
$ kubectl apply -f
```

<https://raw.githubusercontent.com/kubernetes/dashboard/v2.6.1/aio/deploy/recommended.yaml>

```
kubectl get svc --all-namespaces
```

```
kubectl -n kubernetes-dashboard get service kubernetes-dashboard
```

```
kubectl -n kubernetes-dashboard edit service kubernetes-dashboard
```

change ClusterIp to NodePort

```
kubectl -n kubernetes-dashboard get service kubernetes-dashboard
```

```
.
apiVersion: v1
kind: Service
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"labels":{"k8s-app":"kubernetes-dashboard"},"name":"kubernetes-
creationTimestamp: "2019-05-20T14:19:11Z"
labels:
  k8s-app: kubernetes-dashboard
name: kubernetes-dashboard
namespace: kube-system
resourceVersion: "38440"
selfLink: /api/v1/namespaces/kube-system/services/kubernetes-dashboard
uid: 3d042c00-7b0a-11e9-bbe4-0257546da840
spec:
  clusterIP: 10.105.47.198
  ports:
  - port: 443
    protocol: TCP
    targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
  sessionAffinity: None
  type: NodePort
status:
  loadBalancer: {}
```

```
ubuntu@ip-172-31-46-112:~$ kubectl -n kubernetes-dashboard get service kubernetes-dashboard
NAME                                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes-dashboard               NodePort    10.100.75.69   <none>         443:32349/TCP    32m
ubuntu@ip-172-31-46-112:~$
```

take port number allow to security group.

access the dash board using node public ip with nodeport

<https://<nodeip>:<port>/>

it will two options config or else token.

Command line proxy

You can enable access to the Dashboard using the kubectl command-line tool, by running the following command:

```
$ kubectl proxy &
```

Creating a Service Account

We are creating Service Account with name admin-user in namespace kubernetes-dashboard first.

```
ubuntu@ip-172-31-5-3:~$ vi dashboard-adminuser.yaml
```

```
apiVersion: v1
```

```
kind: ServiceAccount
```

```
metadata:
```

```
  name: admin-user
```

```
  namespace: kubernetes-dashboard
```

```
$ kubectl apply -f dashboard-adminuser.yaml
```

Creating a ClusterRoleBinding

In most cases after provisioning cluster using kops, kubeadm or any other popular tool, the ClusterRole cluster-admin already exists in the cluster. We can use it and create only

The screenshot shows the Kubernetes dashboard interface. The top navigation bar includes the 'Workloads' menu. The left sidebar lists various Kubernetes resources: Workloads (Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets), Service (Ingresses, Services), Config and Storage (Config Maps, Persistent Volume Claims, Secrets, Storage Classes), and Cluster (Cluster Role Bindings, Cluster Roles, Events, Namespaces, Network Policies). The main content area displays 'Workload Status' with three green circles representing Deployments, Pods, and Replica Sets, each with a 'Running: 1' indicator. Below this, the 'Deployments' table shows one deployment named 'nginx-server' using the 'nginx:1.19-alpine' image, with labels 'app: nginx' and '1 / 1' pods created '15 hours ago'. The 'Pods' table shows one pod named 'nginx-server-5bbdd555fb-kb5d4' using the 'nginx:1.13-alpine' image, with labels 'app: nginx' and 'pod-template-hash: 5bbdd555fb', running on node 'ip-172-31-5-218' with 1 restart. The 'Replica Sets' section is partially visible at the bottom.

Name	Images	Labels	Pods	Created
nginx-server	nginx:1.19-alpine	app: nginx	1 / 1	15 hours ago

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
nginx-server-5bbdd555fb-kb5d4	nginx:1.13-alpine	app: nginx pod-template-hash: 5bbdd555fb	ip-172-31-5-218	Running	1	-	-	15 hours ago

Remove the admin ServiceAccount and ClusterRoleBinding.

```
kubectl -n kubernetes-dashboard delete serviceaccount admin-user
```

```
kubectl -n kubernetes-dashboard delete clusterrolebinding admin-user
```